

ChamiNext × Apex Labs — POC One-Pager

Joint POC Workshop · Uzziel Perez & Rich Lanot, Founders · June 2026 **Ask:** Design partners + Bedrock pilot (credits, dedicated support, warm SMB intros) **Demo:** [chaminext.netlify.app](#)

The Idea

ChamiNext is the **evaluation layer for engineering in the AI era**: *practice like you work, ship real output, hire on signal — not puzzle count*. The bet: in 2026 the real differentiator is **how well someone reasons and ships alongside AI** — prompting judgment, verification discipline, recovery when AI output is wrong — not recall under a whiteboard.

This workshop pressure-tests that direction and explores **AWS Bedrock** as the infrastructure for an agentic evaluation layer — the engine behind ChamiNext's AI Mocks, Ship Tests, and the employer Interview Studio. Goal: leave with a scoped, shippable POC — not a slide deck.

What We Evaluate

ChamiNext's three product lines (AI Mocks · Ship Tests · Interview Studio) all run on one shared grading engine. Four rubric dimensions:

Dimension	What it tests	Product line	Scope
Reasoning (Adaptive AI)	Socratic reasoning, decomposition, AI-assisted problem solving — not recall	AI Mocks	MVP
Ship test	Build + deploy a real feature in a 24–72h window; code quality + delivery judgment	Ship Tests	MVP
System design	Architecture trade-off reasoning under ambiguity	Studio	Phase 2
AI collaboration vibe check	Live prompt specificity, trust calibration, recovery	All	Phase 2

POC = the shared grading engine, validated first on the two most ground-truth-able tasks (reasoning-graded coding task + ship test) as an **LLM-as-judge over a static submission**. Phases 3–4 add live-session capture.

POC Architecture

Submission ingest → Rubric structuring → LLM-as-judge eval → Score + cited rationale

- **S3** — submission store (repo / zip / PR link)
- **Bedrock (Claude 3.5 Sonnet)** — fixed rubric per task type; one batched judge call per submission
- **Bedrock Guardrails** — bound outputs, catch hallucinated code claims
- **Cognito + RDS** — identity and structured score/telemetry storage (per deck pilot architecture)
- **Cost target** — €0.05–€0.30 per eval; single judge call, rationale capped ~300 tokens (no iterative loop in POC)

Scope

In: one reasoning-graded coding task + one ship-test task · fixed rubrics · score + per-dimension breakdown + cited rationale · reviewer validation view · validation against labeled human-grader set

Out: candidate-facing UI · scheduling · multi-language generalization · cheating detection · live-session capture · system design / vibe-check pillars

Timeline — 6 weeks + 3 week buffer

This is the engineering core of the deck's **90-day collaborative Bedrock pilot** (build in weeks 1–9; remaining time for design-partner validation and GTM co-sell).

Sprint	Week	Goal
0 — Workshop & Scope	0	Align architecture, model, scope, credits path
1 — Foundation	1	AWS + Bedrock access, S3 ingest, skeleton pipeline
2 — Reasoning judge	2	First working eval, rubric v1, score + rationale
3 — Ship-test judge	3	Extend pipeline, code-line citations
4 — Validation	4	Agent vs. human grader, variance check, Guardrails
5 — Reviewer view	5	Auditable rationale, outlier flagging
6 — Harden & cost	6	Cost/latency report, demo-ready POC
Buffer	7–9	Absorb slippage, iterate on model/rubric
Stretch (if ahead)	7–9	Conversational AI Mock prototype — proves the engine generalizes to the deck's flagship experience

Milestones: M1 (Wk1) ingest end-to-end · M2 (Wk3) both pillars scored + cited · M3 (Wk4) scores correlate with human set · M4 (Wk6) demo-ready with cost/latency numbers.

Why This Matters to Hiring Teams

The deck's own pain metrics are the bar this POC must clear:

- **Cuts false positives/negatives** — whiteboard loops yield high error rates; the POC must show its scores correlate with human graders.
- **Portable, verifiable proof** — a deployed build + cited rationale a recruiter trusts without re-reading the submission.
- **AI-cheating clarity** — distinguishes genuine AI-assisted development from copy-paste, by grading reasoning and process, not just output.
- **Frees senior engineer time** — replaces hours of whiteboard theater per loop with auto-ranked, auditable signal.

Open Questions for Apex Labs

1. **Model selection** — most consistent, auditable Bedrock model at this complexity? Variance across repeated runs?
2. **Guardrails vs. custom harness** — best fit for bounding scores and catching hallucinated code claims?
3. **Rationale traceability** — tool use, structured output, or post-processing for code-line citations?
4. **Human-in-the-loop** — existing AWS patterns for the reviewer handoff?
5. **Cost per eval** — Bedrock cost at this token budget, scaled to hundreds of submissions per hiring cycle?

Success Criteria

Workshop: concrete Bedrock model + architecture direction · agreed POC scope and next steps · clarity on credits, dedicated support, and warm SMB intros (co-sell path).

POC: agent scores correlate strongly with human graders · rationale specific enough to trust without re-reading the submission · runs end-to-end on Bedrock at acceptable latency and €0.05–€0.30 per eval.